

Latent Dimensionality Reduction: An Algorithmic Method for Interpreting Black Box Models

DRAFT

Elias Kassell, Fred Farrell

September 2019

1 Abstract

LDR (Latent Dimensionality Reduction) is an algorithmic method for interpreting models by decomposing predictions from the original high dimensional feature space into a lower dimensional space. Both the output of the model and the certainty of its predictions are encoded in the condensed space. The method works out of the box on both classification and regression problems with input features relating to interpretable metrics. The output gives the ability to glimpse higher dimensional spaces, which are normally too complex for humans to understand, and is a powerful tool in understanding the underlying features generating the metrics.

Acknowledgements

Thanks to Illumina® for supporting this research. The extremely high dimensional problems posed by genomic analysis provided the inspiration for this work.

2 Introduction

When training a model to make predictions in situations where interpretability is not required, the main goal is to maximise predictive accuracy [1]. Often black box models have a higher predictive accuracy than more interpretable solutions, particularly in higher dimensional or more complex problems [2]. Because of the uninterpretability of these models, they become impractical for situations where it is a necessity to understand why a model has chosen a particular output as its prediction, and what effect changes to values of the input features will have on the output of the model.

Some examples of uses of model interpretation in real world applications include

1. Understanding the quality of a systems current understanding [3].
2. Interpreting how the value of a feature, or subset of features, affects a model's prediction (henceforth referred to as "feature interpretation"). For example, in a medical setting this interpretation can be used for suggesting the relation between features

and a diagnosis. In another example, for manufacturing this interpretation can provide target values for the underlying processes involved in quality of produce.

3. The ability to use a model when not all values for the input features are present. By integrating across features which are not present, the certainty of a models prediction can be estimated. Current methods, such as imputing values for features which are not present, ignores the variability of the absent features across the feature space, which LDR can provide.

3 The Algorithm

1. Scale the original data so it falls into $[0, 1]$ intervals.
2. Train a predictive model on the scaled data.
3. Train an outlier detection predictive model on the scaled data.
4. Create a KDE (kernel density estimate) [4] of the training samples with a bandwidth h (a good value for this is $\frac{3}{r}$, as this spreads the points out more). Optional: to improve the visualization, normalize the training to produce a constant density output.
5. Sample n new points from the kernel density, using both classifiers to make predictions of each point. Weight the prediction according to the outlier classifier, where outliers have no model certainty in the prediction at that point.
6. Bin the samples according to regular intervals. For each dimension, group points with resolution r , reducing the value of the bin to the mean prediction across it.

4 Classification Example

The Wisconsin breast cancer dataset [5] (provided by scikit [6]) is used as it has 31 dimensions, a significant number. The aim with the dataset is to classify the tumor as either malignant or benign. There are 569 samples in total, with 357 benign and 212 malignant.




Figure 1: RF prediction certainty, where red indicates certainty of malignancy, blue certainty of benignity, and yellow a general lack of certainty. A lack of drawn contour indicates no samples were drawn from the sample space, as the kernel density was too low in those areas. The black dots are the locations of the samples. The top and left axis are the single dimensional decomposition, while the bottom right is the two dimensional decomposition. In the single dimension estimates, the resulting value is shifted down by 0.5 to create a flat midpoint for the bars. The axis are labelled with the relative input minimum and maximum values used for scaling.




Figure 2: IF prediction certainty, where red indicates certainty of outlier, blue certainty of not outlier, and yellow uncertainty.

4.1 Random Forest and Isolation Forest Interpolation

The 100 estimator RF (Random Forest) [7] achieved an F1 score of 0.975, while an IF (Isolation Forest) [8] with a dynamically selected frontier was trained for outlier detection. Benign was selected as the positive category, resulting in a value of 1.0 indicating benign tumors and 0.0 indicating malignant. All of the data is min-max scaled, then a 0.7/0.3 training/testing data split is applied. LDR is applied to the training data, integrating over 50,000 randomly selected points from the weighted kernel density of the sample space with a resolution of 16.66. Mean symmetry and mean area are selected for visual inspection of their individual effect on the model’s classification. The certainty of the RF is shown in figure 1, while the certainty of the IF is shown in figure 2. The final combined certainty is shown in figure 3.

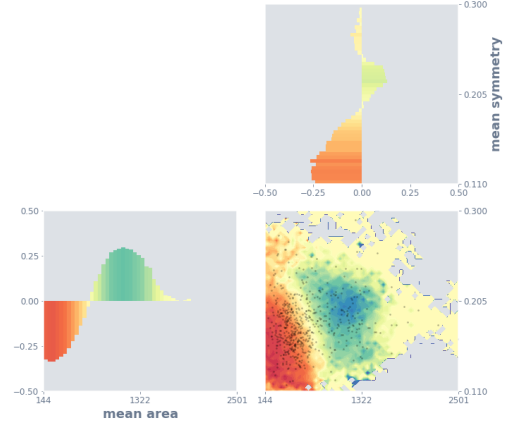


Figure 3: RF with IF interpolation prediction certainty, where red indicates overall certainty of malignancy, blue indicates overall certainty of benignity, and yellow general overall uncertainty.




Figure 4: NN with IF interpolation prediction certainty.

4.2 Neural Network and IF Interpolation

A 3 layer feed forward Neural Network (NN) was used (via pytorch [9]), consisting of an input layer, a hidden layer where tanh is applied, and an output layer. Soft-max is applied when making predictions, and the final certainty of the prediction is calculated by calculating the proportion of the largest weighted class out of all weightings. The NN was trained over 50,000 epochs and achieved an F1 score of 0.961, slightly worse than the RF. The resulting LDR of the same features can be seen in figure 4

4.3 Analysis of Discrepancy Between RF and NN

There is very little discrepancy between the two LDR representations, both of which have similar predictive success. This is significant because random forests are widely regarded as more interpretable than neural networks [10], but LDR represents them both in an equally interpretable fashion. The only visual difference between the two LDR representation arises from the NN tending to be more certain of its classifications. This




Figure 5: RF prediction certainty scatter plot of mean symmetry.




Figure 6: NN prediction certainty scatter plot of mean symmetry.

can be seen in the scatter plots of prediction certainties for mean symmetry of the RF in figure 5m compared to the NN in figure 6.

5 Regression Example

A similar method as with the classification was applied to the Boston House Price dataset [11] (also provided by scikit). This dataset has 14 dimensions, where the target is to predict the house price based off of quantitative factors such as age of the house, the number of rooms, and the per capita crime rate of the time within which is situated. Some interesting relations are made immediately available from the LDR, shown in figure 7, are

1. An increasing crime rate tends to correlated with houses being worth less. Above the median level of crime however, the model becomes uncertain.
2. Any quantity of crime only has a negative effect on house price.
3. The lower the crime rate, the more rooms a house tends to have, and the more it is worth.
4. The older the house, the less it tends to be worth.
5. The age of the house and the number of rooms are for the most part independent.

The prediction certainty prior to be binning can be seen in figure 8.

6 Mathematical Explanation

6.1 Feature Space and Subsets Definition

Let a scaled feature space Ω for samples $w = \{0, 1, \dots, l\}$ be described as

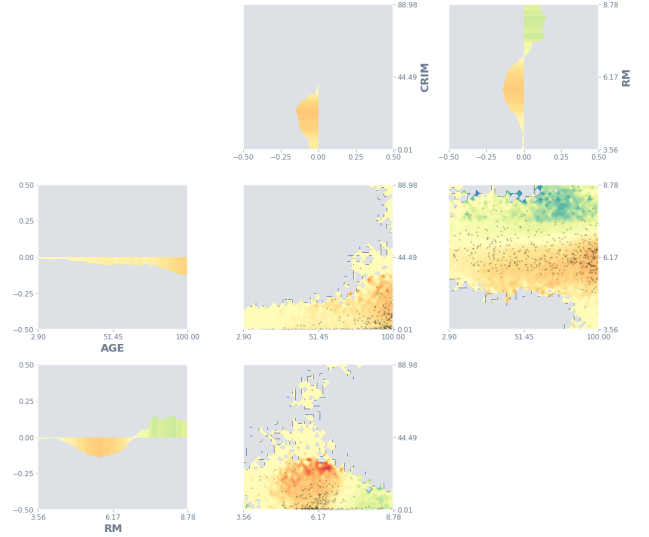


Figure 7: Regression RF with IF interpolation. AGE is the age of the house, RM the number of rooms, and CRIM the per capita crime rate per town.




Figure 8: RF prediction certainty of number of rooms.

$$\Omega = \begin{bmatrix} \omega_1 = \{d_{1,1}, d_{2,1}, \dots, d_{l,1}\} \\ \omega_2 = \{d_{1,2}, d_{2,2}, \dots, d_{l,2}\} \\ \vdots \\ \omega_m = \{d_{1,m}, d_{2,m}, \dots, d_{l,m}\} \end{bmatrix}$$

where $d_{i,j} \in [0, 1]$, l is the number of features and m is the number of samples in the feature space.

The subset X of values for features $t \subset w$ which are being used to predict a set Y of features $l - t = \{i, (i \in w) \wedge (i \notin t)\}$ are therefore defined as

$$X = \begin{bmatrix} x_1 = \{d_{1,1}, d_{2,1}, \dots, d_{|t|,1}\} \\ x_2 = \{d_{1,2}, d_{2,2}, \dots, d_{|t|,2}\} \\ \vdots \\ x_m = \{d_{1,m}, d_{2,m}, \dots, d_{|t|,m}\} \end{bmatrix}$$

$$Y = \begin{bmatrix} y_1 = \{d_{1,1}, d_{2,1}, \dots, d_{|l-t|,1}\} \\ y_2 = \{d_{1,2}, d_{2,2}, \dots, d_{|l-t|,2}\} \\ \vdots \\ y_m = \{d_{1,m}, d_{2,m}, \dots, d_{|l-t|,m}\} \end{bmatrix}$$

6.2 Model and Outlier Prediction Interpolation

Let p describe a model prediction mapping, such that $p(x_i) \approx y_i$. Let o describe an outlier prediction mapping such that $o(x_i) = [1, 0]$, where 0 indicates complete certainty of x_i being an outlier and 1 complete

certainty of x_i not being an outlier. Let outlier weighting and model prediction interpolation function f be defined as

$$f(x_i) = (p(x_i) - 0.5) \times o(x_i) + 0.5.$$

6.3 KDE Sample Drawing

The KDE with bandwidth $1/r$ can be estimated and n samples drawn from it in unison by drawing from the individual weighting each sample $x \in X$ contributes to the kernel density. This is done by randomly selecting n samples from X and drawing distributed points $q \sim N(x, 1/r)$. This sampling is equivalent to the VEGAS algorithm for the subsequent monte carlo integration [12]. The set of samples Q can therefore be described as

$$Q = \{q_1, q_2, \dots, q_n\}.$$

Note that the VEGAS algorithm relies on the assumption that the sample space is univariate and i.i.d, resulting in LDR only successfully interpreting samples from the same feature space.

6.4 Monte Carlo Integration Over Feature Subsets

Let e be the set of intervals describing a single dimensional set of bins, such that

$$e = \{[0, \frac{1}{r}], (\frac{1}{r}, \frac{2}{r}], \dots, (\frac{r-1}{r}, 1]\}.$$

Let $b \subset t$ describe the feature subspace being interpreted. Let E describe the $|b|$ multivariate permutations of bin pairs, such that

$$E = \text{perm}(e_1, e_2, \dots, e_{|b|}).$$

The Monte Carlo integration of a single point of the feature subset can therefore be described, where Q_k represents all elements that fall within the bin E_k , as

$$I_{E_k} = \frac{1}{z} \sum_{i=1}^z f(q_i), q_i \in Q_k.$$

The entire feature subset can be interpreted by applying this for all $E_K \in k$.

7 Conclusion

I have demonstrated that it is possible to interpret a models understanding of high dimensional features spaces by estimating the certainty of the model over lower dimensional subsets of the feature space. However, before acceptance for clinical purposes, the model should be rigorously tested on more data sets. In addition to this, further research of extending LDR could be to interpret latent features that do not relate to uninterpretable input features, such as neural networks trained on images; a visualization of the prediction of the model based on a single pixel would not be of much use by itself.

References

- [1] Leo Breiman et al. “Statistical modeling: The two cultures (with comments and a rejoinder by the author)”. In: *Statistical science* 16.3 (2001), pp. 199–231.
- [2] Jürgen Schmidhuber. “Deep learning in neural networks: An overview”. In: *Neural networks* 61 (2015), pp. 85–117.
- [3] Saleema Amershi et al. “Effective end-user interaction with machine learning”. In: *Twenty-Fifth AAAI Conference on Artificial Intelligence*. 2011.
- [4] Emanuel Parzen. “On estimation of a probability density function and mode”. In: *The annals of mathematical statistics* 33.3 (1962), pp. 1065–1076.
- [5] Olvi L. Mangasarian Dr. William H. Wolberg W. Nick Street. *Breast Cancer Wisconsin (Diagnostic) Data Set*. Nov. 1995. DOI: [https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Diagnostic\)](https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+(Diagnostic)).
- [6] Fabian Pedregosa et al. “Scikit-learn: Machine learning in Python”. In: *Journal of machine learning research* 12.Oct (2011), pp. 2825–2830.
- [7] Andy Liaw, Matthew Wiener, et al. “Classification and regression by randomForest”. In: *R news* 2.3 (2002), pp. 18–22.
- [8] Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. “Isolation forest”. In: *2008 Eighth IEEE International Conference on Data Mining*. IEEE. 2008, pp. 413–422.
- [9] Adam Paszke et al. “Automatic Differentiation in PyTorch”. In: *NIPS Autodiff Workshop*. 2017.
- [10] Lin Song, Peter Langfelder, and Steve Horvath. “Random generalized linear model: a highly accurate and interpretable ensemble predictor”. In: *BMC bioinformatics* 14.1 (2013), p. 5.
- [11] D. Harrison and D.L. Rubinfield. *The Boston house-price data*. Nov. 1978. DOI: <http://lib.stat.cmu.edu/datasets/boston>.
- [12] G Peter Lepage. “A new algorithm for adaptive multidimensional integration”. In: *Journal of Computational Physics* 27.2 (1978), pp. 192–203.